

The background features a complex network of thin grey lines connecting various points, creating a web-like structure. Scattered throughout are numerous triangles of different sizes and orientations, some solid and some outlined. The overall aesthetic is modern and technical.

Team 10

Pricing and Advertising

Daglio Gabriele, Di Cesare Federico, Germano Jacopo
23th July 2021

The background features a complex network of thin, light gray lines connecting various dark gray circular nodes. These nodes are distributed across the frame, with a higher concentration in the upper right and lower right areas, creating a web-like or molecular structure. The overall aesthetic is minimalist and technical.

Environment

The users

Features: F_1 and F_2 , independent.

$$\Pr(F_1 = \textit{True}) = \theta_1$$

$$\Pr(F_2 = \textit{True}) = \theta_2.$$

Likelihoods:

$$\tilde{l}_{TT} = \theta_1 \theta_2$$


$$\tilde{l}_{TF} = \theta_1 (1 - \theta_2)$$

$$\tilde{l}_{FT} = (1 - \theta_1) \theta_2$$

$$\tilde{l}_{FF} = (1 - \theta_1)(1 - \theta_2)$$



Example



One round

λ_a = Average number of daily auctions.

$v(b)$ = Probability of winning an auction with bid b

$r_{comb}(p)$ = Conversion rate of *comb* users for price p

f_{comb} = Average future visits of *comb* users

Each day:

1. $A \sim \text{Poisson}(\lambda_a)$ # number of auctions
2. $A_{TT}, A_{TF}, A_{FT}, A_{FF} \sim \text{Multinomial}(A, (\tilde{l}_{TT}, \tilde{l}_{TF}, \tilde{l}_{FT}, \tilde{l}_{FF}))$
3. $N_{comb} \sim \text{Binomial}(A_{comb}, v(b))$ # the new clicks
4. $C_{comb,j} \sim \text{CostPerClick}(k_{comb}(b))$ # cost of click j
5. $D_{comb,i} \sim \text{Binomial}(N_{comb}, r_{comb}(p))$ # whether user i purchases
6. $F_{comb,i} \sim \text{Poisson}(f_{comb})$ # future visits user i



Step 1

Offline algorithm

Without context generation

Expected profit

$$ExpectedProfit(p, b) = E \left[\sum_{c \in C} \sum_{i=1}^{N_{c,b}} \left(D_{c,p,i} (1 + F_{c,i}) m(p) - C_{c,b,i} \right) \right]$$

\Downarrow

$$ExpectedProfit(p, b) = \lambda_a v(b) \sum_{c \in C} l_c \left(r(c, p) (1 + f(c)) m(p) - k(c, b) \right)$$


Independent Price Hierarchy

$p_1 \succsim p_2$ for some bid



$p_1 \succsim p_2$ for every bid



The algorithm

1. $\bar{b} \leftarrow \text{median}(B)$
2. $p^* \leftarrow \text{optimal_price_fixed_bid}(\bar{b})$ # linear scan of the prices
3. $b^* \leftarrow \text{optimal_bid_fixed_price}(p^*)$ # linear scan of the bids
4. *return* (p^*, b^*)



Step 2

Online problem formalization



Online setting

- The averages are not know
- The feedback is partially delayed
 - The future visits come with a delay of 30 rounds.



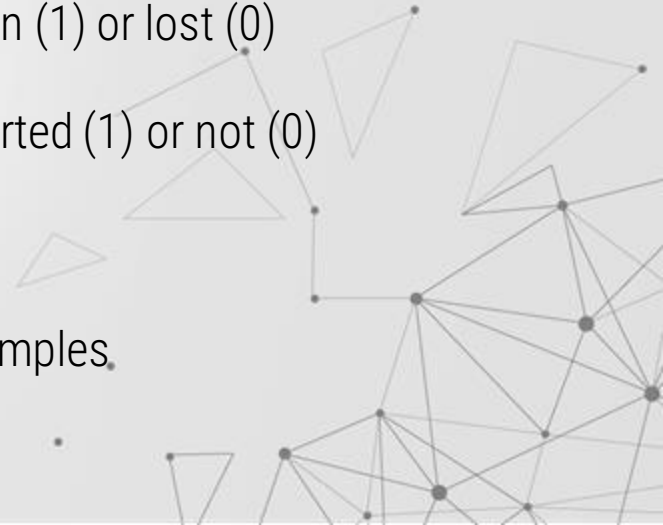
Online setting - Feedback

- After round i , the learner receives:
 - The auctions run on round i
 - The new clicks obtained on round i
 - The purchases performed on round i
 - The cost paid for the clicks of round i
 - The future visits of the purchases of round $i - 30$



Online setting - Estimates

- The averages are not known
- Need to estimate:
 - Auctions --> Sample mean
 - New clicks --> Bernoulli variable, auction won (1) or lost (0)
 - Purchases --> Bernoulli variable, click converted (1) or not (0)
 - Cost per click --> Sample mean
 - Future visits --> Sample mean of complete samples.



Estimate the profit

$$\widehat{exp(p, b)}_i = \bar{a}_i \bar{w}_{bi} \left(m(p) \bar{r}_{pi} (1 + \bar{f}_{pi}) - \bar{c}_b \right)$$

Optimistically explore w and r



Optimistic exploration

Two approaches:

Upper Confidence Bound

- Conversion rate --> Conversion rate upper bound
- Bid win rate --> Bid win rate upper bound

Thompson Sampling

- Conversion rate --> Sample from Betas, update the Betas
- Bid win rate --> Sample from Betas, update the Betas





Step 3

Online pricing

Without discriminating between customer classes

OUR NUMBERS

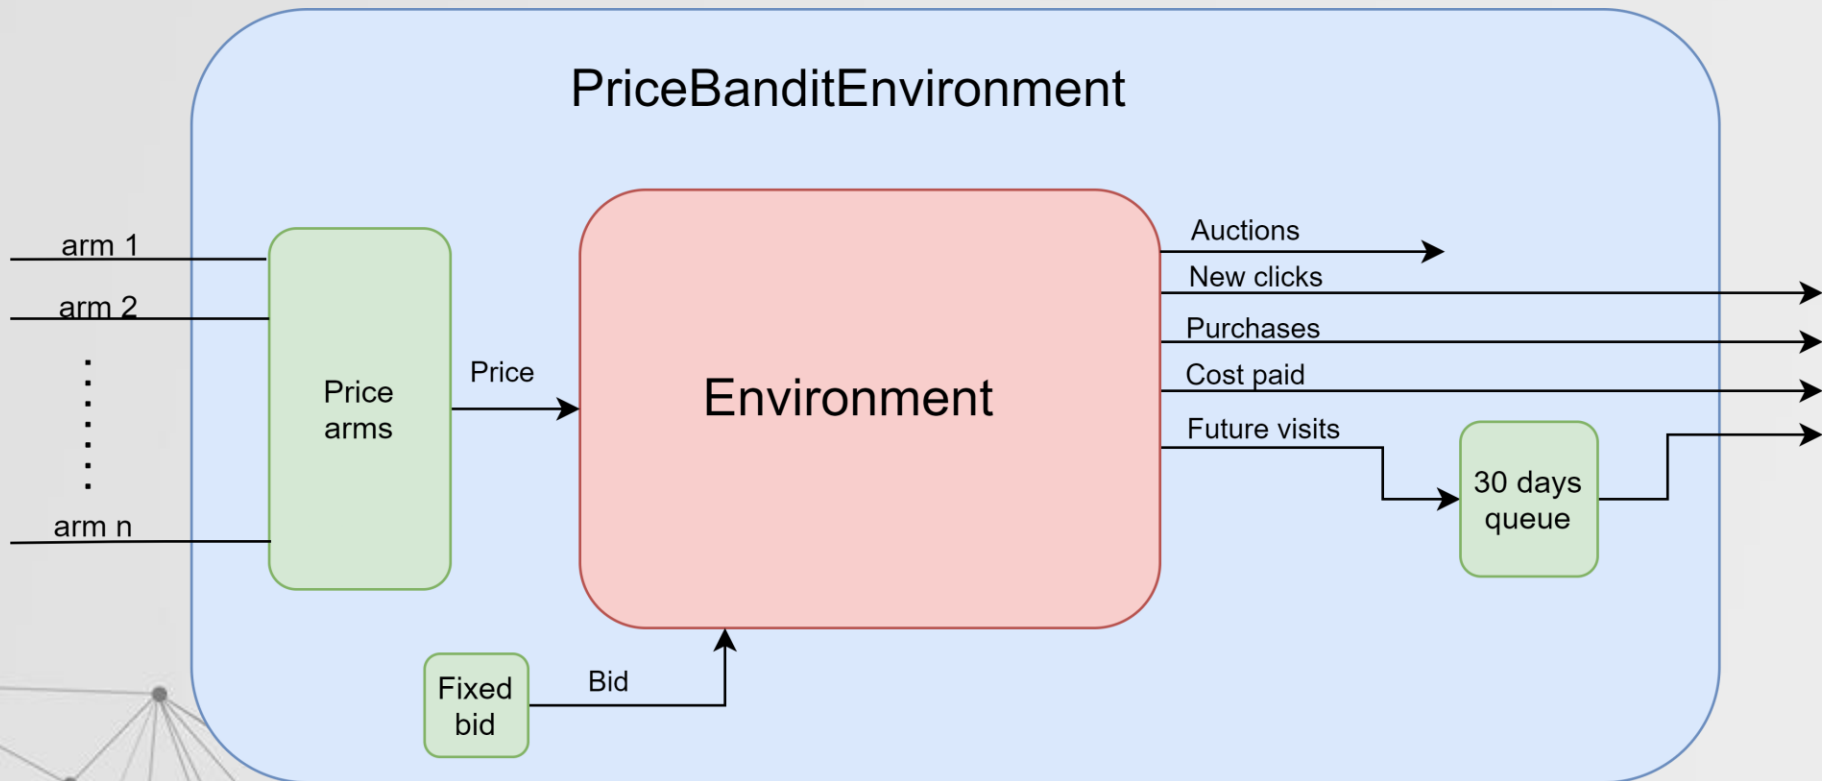
10

PRICES

1

BID





PriceBanditEnvironment

Purchases

New clicks

Cost paid

Future visits

OptimalPriceLearner

Pull from env

Choose next arm

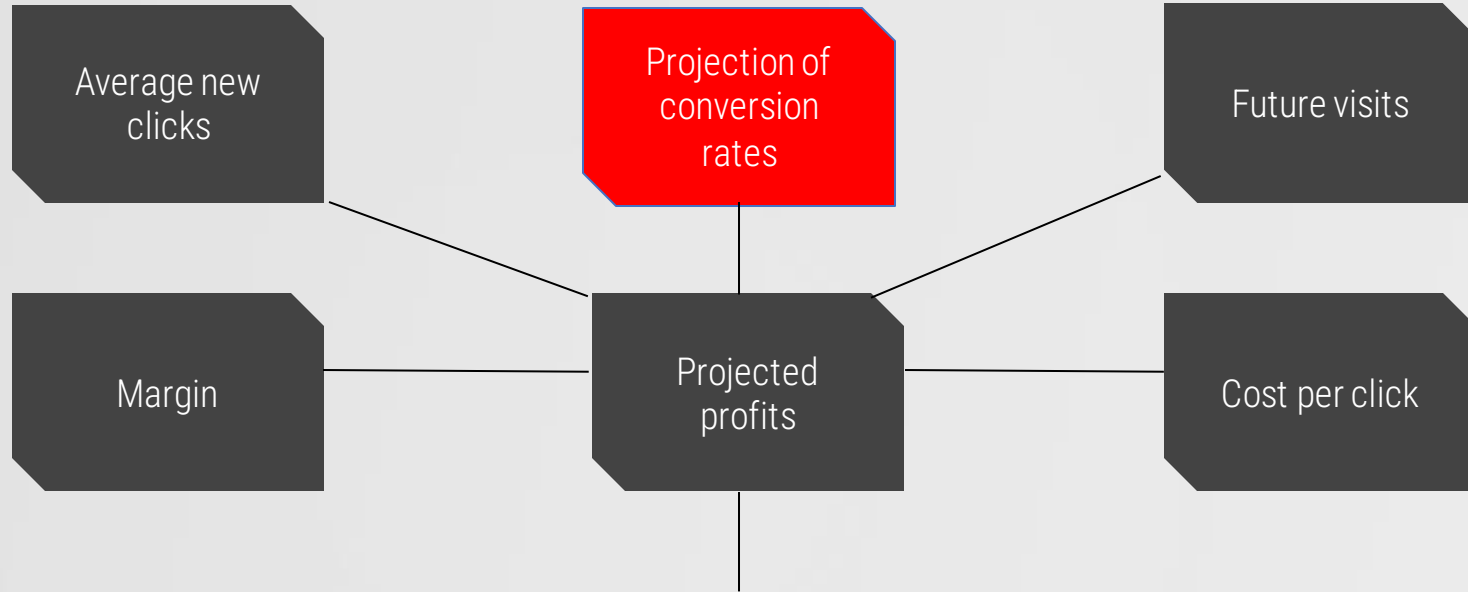


INITIAL ROUND-ROBIN

- Round robin for $30 + \textit{number of arms}$ rounds
- Until 1 full feedback for each arm



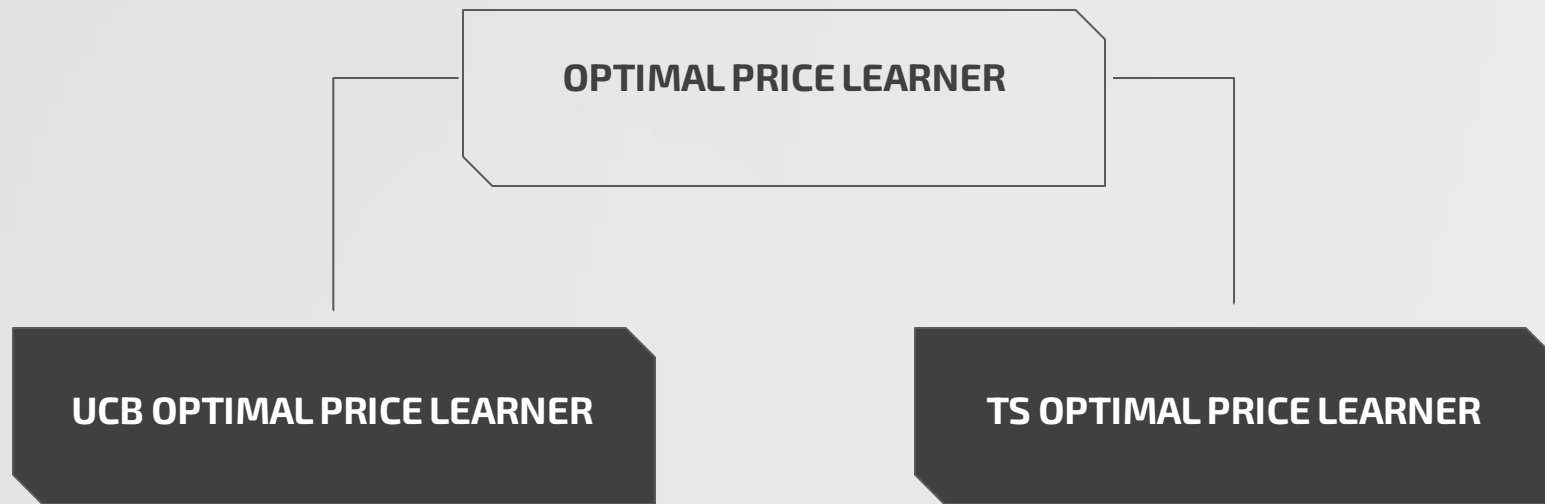
CHOICE OF THE NEXT ARM




NEXT ARM \leftarrow ARGMAX



PROJECTION OF CONVERSION RATES



UCB OPTIMAL PRICE LEARNER



UCB OPTIMAL PRICE LEARNER

$$\frac{\textit{purchases}_a}{\textit{clicks}_a} + \sqrt{\frac{2\log(t)}{\textit{clicks}_a}}$$

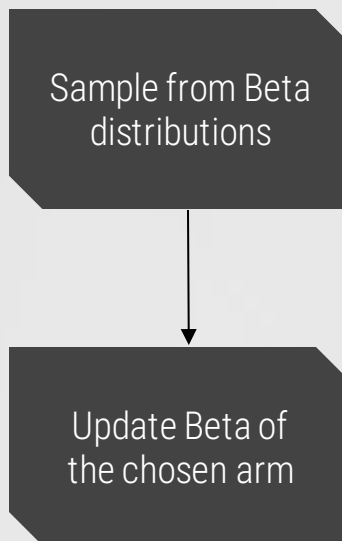
AVERAGES

RADII

The updates happen in **batches** since, at each round, we see the realization of **new_clicks** tries








TS OPTIMAL PRICE LEARNER



The updates of the parameters α and β of the Beta distributions happen in **batches** since, at each round, we see the realization of ***new_clicks*** tries

Seed: 3144548588









Seed: 3144548588


Price	Expected	Norm. gaps	UCB Pulls	UCB Expected	TS Pulls	TS Expected
10.00	-4309.67	100.0	4	-4287.03	4	-4447.41
20.00	-48.64	43.9	4	16.98	4	28.71
30.00	2460.29	10.9	10	2650.07	4	2066.38
40.00	3291.72	0.0	141	3266.95	302	3344.78
50.00	2977.05	4.1	131	3045.23	31	3016.02
60.00	1924.81	18.0	34	2020.80	4	1859.43
70.00	632.38	35.0	13	218.62	4	-180.93
80.00	-440.08	49.1	13	-435.83	4	-165.58
90.00	-1144.79	58.4	11	-1236.55	4	-1492.98
100.00	-1550.37	63.7	4	-1806.97	4	-1715.24

Optimal price: 40.00, Optimal bid: 27.00



Seed: 1873674269





Seed: 1873674269

Price	Expected	Norm. gaps	UCB Pulls	UCB Expected	TS Pulls	TS Expected
10.00	-1564.32	100.0	4	-1677.11	4	-1704.78
20.00	-806.12	73.9	4	-775.53	4	-778.56
30.00	-172.82	52.1	4	-200.15	4	-153.46
40.00	361.10	33.8	4	264.64	4	414.62
50.00	808.49	18.4	5	765.44	4	993.70
60.00	1152.60	6.5	14	1044.73	27	1177.04
70.00	1342.60	0.0	95	1351.48	261	1351.32
80.00	1327.18	0.5	122	1336.60	7	825.18
90.00	1131.00	7.3	65	1138.87	45	1097.34
100.00	865.62	16.4	48	967.43	5	666.85

Optimal price: 70.00, Optimal bid: 9.00





Step 4

Online pricing

Discriminating between customer classes

OUR NUMBERS

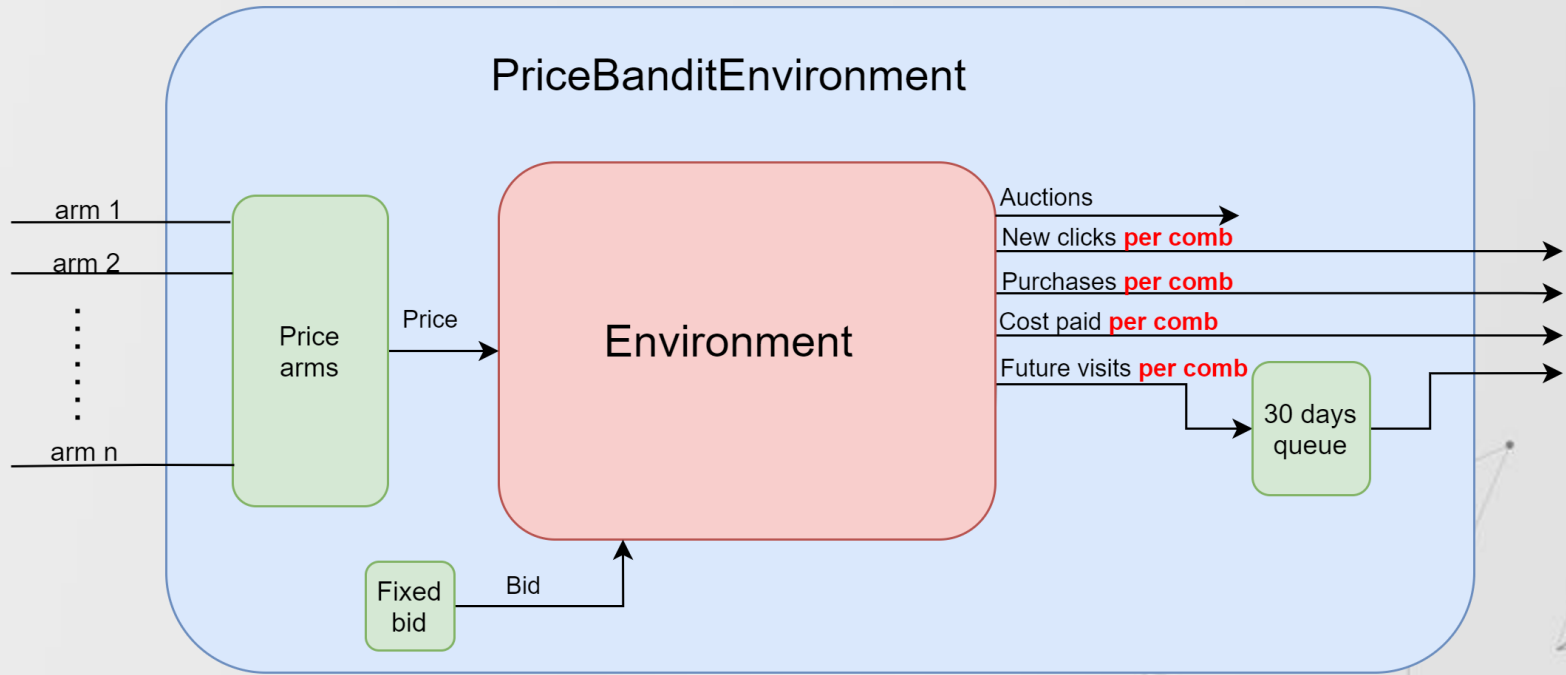
10

PRICES

1

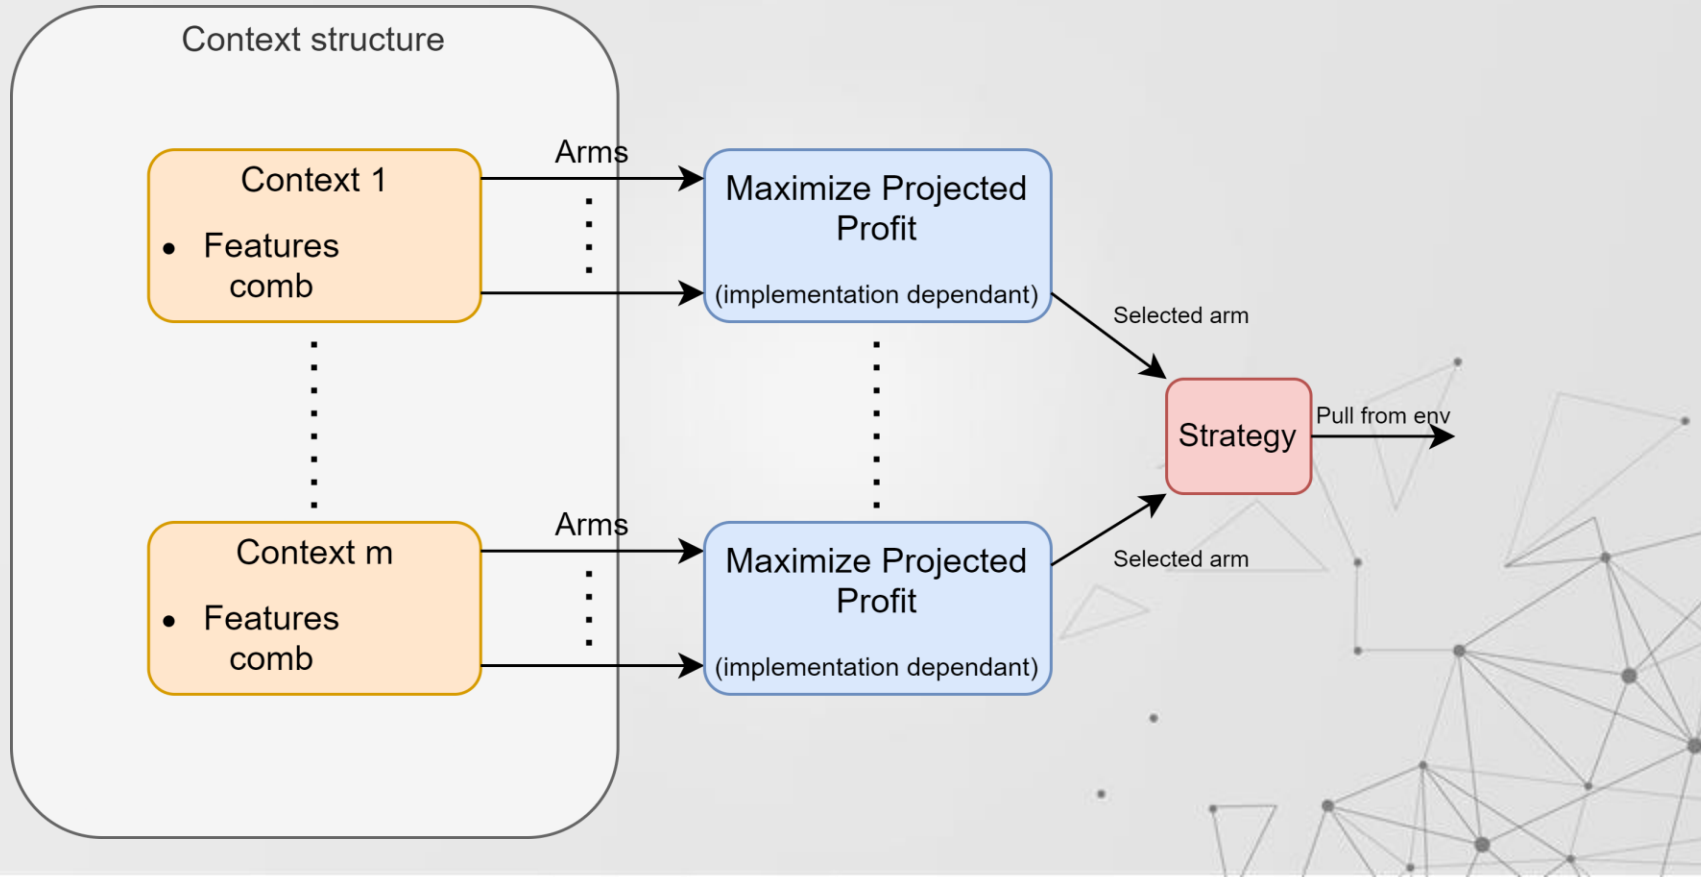
BID





- Now, PriceBanditEnvironment returns new clicks, purchases, cost paid and future visits **per comb**

OPTIMAL PRICE DISCRIMINATING LEARNER



CHOICE OF THE NEXT STRATEGY

- Just like Step 3
- UCB1 and Thompson Sampling are applied **context-wise**



UCB OPTIMAL PRICE DISCRIMINATING LEARNER

$$\frac{\textit{purchases}_a}{\textit{clicks}_a} + \sqrt{\frac{2\log(t)}{\textit{clicks}_a}}$$

AVERAGES

RADII



TS OPTIMAL PRICE DISCRIMINATING LEARNER

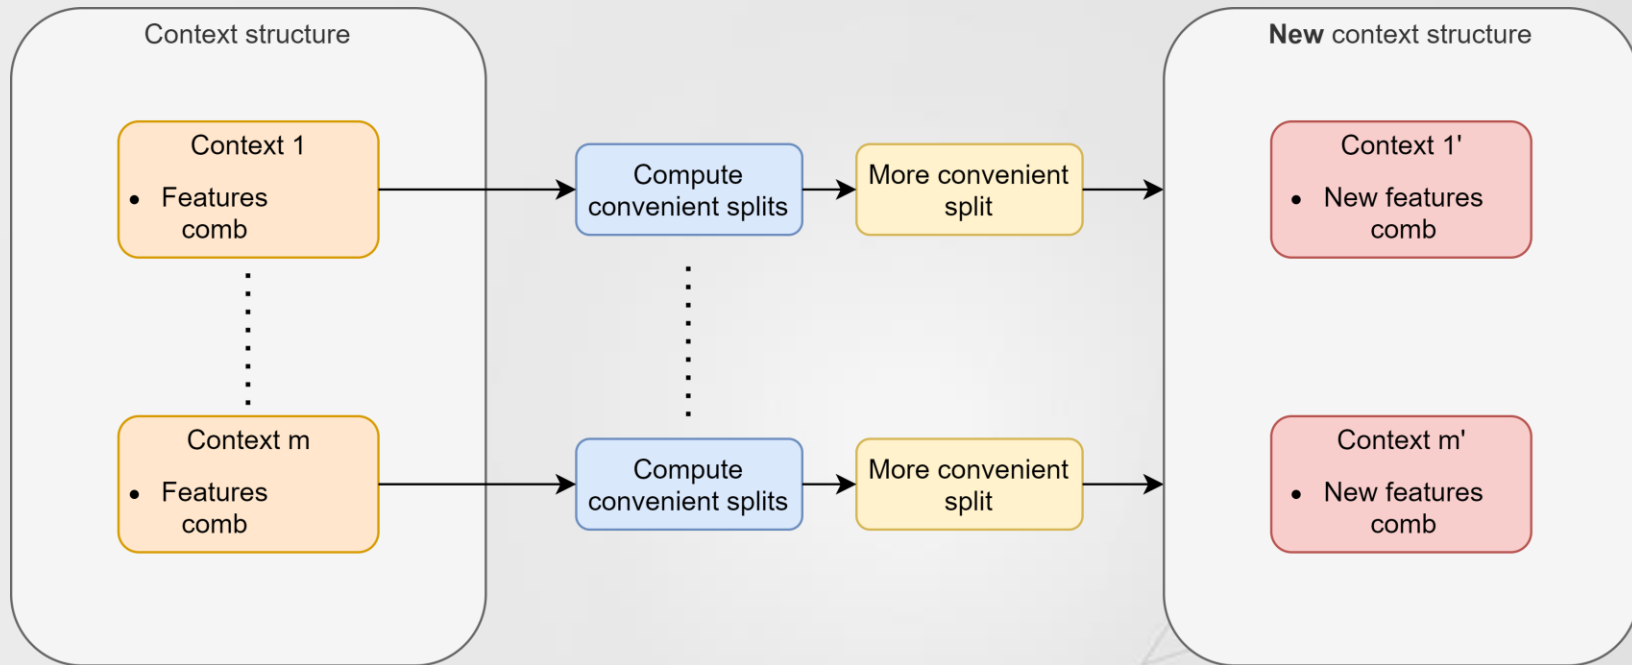
Sample from Beta
distributions



Update Beta of
the chosen arm



CONTEXT STRUCTURE UPDATE



- For each context, compute the **convenient splits** (the lower bound of expected profit after the split is **higher** than the current lower bound)
- For each context, choose the **more convenient** split and create a new context structure

EXPLORATION

- Splits that could not be detected after the initial round robin are never detected, if their optimal arm is very sub-optimal for the current context
- Lower confidence bound never gets tighter



EXPLORATION

- **Explorative rounds:** contexts that could split perform round-robin selection
- Tightens the bounds of the current sub-optimal arms







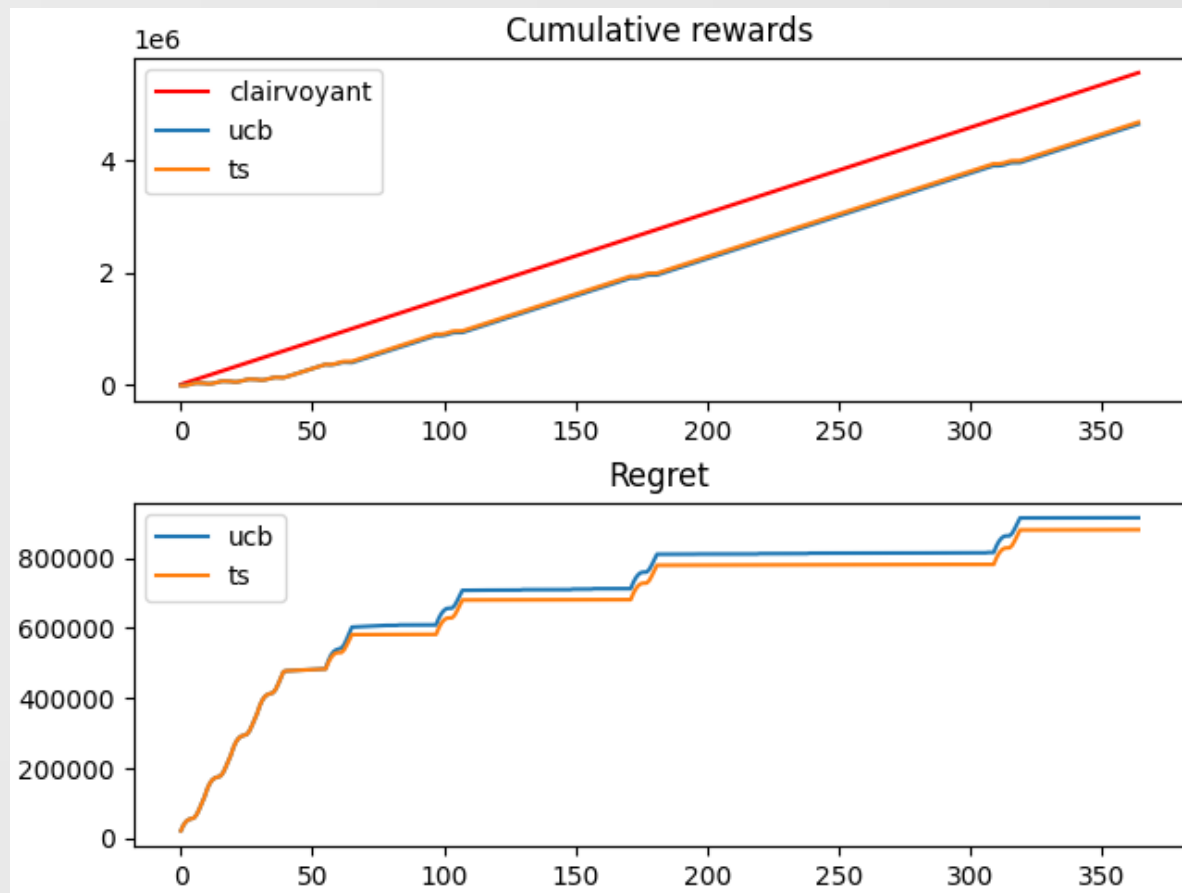
HOW MANY EXPLORATIVE ROUNDS?

- Must be logarithmic
- After the completion of i^{th} explorative round, 2^i normal rounds must be completed before the next explorative round



Seed: 3511939391





Seed: 3511939391

Legend: context, normalized gap, number of pulls

Optimal pricing strategy: C1(TF, FF): 60.00, C2(FT): 90.00, C3(TT): 40.00

UCB with context generation:

Price	TF,FF	Gaps	Pulls	TT	Gaps	Pulls	FT	Gaps	Pulls
10.00		100.0	8		100.0	5		100.0	5
20.00		75.8	8		53.7	5		78.9	5
30.00		51.6	8		11.5	8		59.2	5
40.00		27.6	8		0.0	315		41.4	32
50.00		5.8	14		40.8	6		26.1	5
60.00		0.0	287		65.5	6		13.9	8
70.00		39.1	8		71.0	5		5.4	42
80.00		74.8	8		71.9	5		0.7	43
90.00		83.7	8		72.1	5		0.0	143
100.00		85.2	8		72.1	5		2.8	77

Performed splits:

Round 41 split on feature 2 with incentive 1147.21

Round 78 split on feature 1 with incentive 32.89

Optimal pricing strategy: C1(TF, FF): 60.00, C2(FT): 90.00, C3(TT): 40.00

TS with context generation:

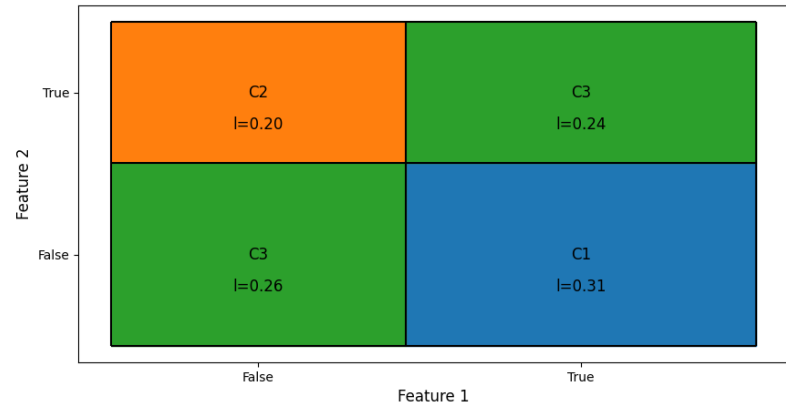



Price	TF,FF	Gaps	Pulls	TT	Gaps	Pulls	FT	Gaps	Pulls
10.00		100.0	8		100.0	4		100.0	4
20.00		75.8	8		53.7	4		78.9	4
30.00		51.6	8		11.5	4		59.2	4
40.00		27.6	8		0.0	328		41.4	17
50.00		5.8	8		40.8	4		26.1	4
60.00		0.0	293		65.5	5		13.9	5
70.00		39.1	8		71.0	4		5.4	17
80.00		74.8	8		71.9	4		0.7	99
90.00		83.7	8		72.1	4		0.0	53
100.00		85.2	8		72.1	4		2.8	158


Performed splits:

Round 41 split on feature 2 with incentive 2994.85

Round 54 split on feature 1 with incentive 122.42

Seed: 1740212098





Seed: 1740212098

Optimal pricing strategy: C1(TF): 30.00, C2(FT): 70.00, C3(TT, FF): 20.00

UCB with context generation:

Price	TF	Gaps	Pulls	FF	Gaps	Pulls	TT	Gaps	Pulls	FT	Gaps	Pulls
10.00		100.0	4		100.0	4		100.0	4		100.0	4
20.00		32.1	4		0.0	4		0.0	6		71.3	4
30.00		0.0	297		43.7	4		43.7	12		46.5	4
40.00		29.3	17		58.6	7		58.6	10		26.5	4
50.00		68.2	6		61.8	125		61.8	20		12.0	12
60.00		85.3	7		62.3	23		62.3	30		3.2	30
70.00		90.6	4		62.4	31		62.4	44		0.0	91
80.00		92.1	18		62.5	43		62.5	59		1.7	89
90.00		92.5	4		62.5	54		62.5	79		7.3	100
100.00		92.6	4		62.5	70		62.5	101		15.5	27

Performed splits:

Round 41 split on feature 2 with incentive 2983.76

Round 42 split on feature 1 with incentive 149.32

Round 44 split on feature 1 with incentive 434.91

Optimal pricing strategy: C1(TF): 30.00, C2(FT): 70.00, C3(TT, FF): 20.00

TS with context generation:

Price	TT	Gaps	Pulls	FT	Gaps	Pulls	TF	Gaps	Pulls	FF	Gaps	Pulls
10.00		100.0	4		100.0	4		100.0	4		100.0	4
20.00		0.0	263		71.3	4		32.1	4		0.0	262
30.00		43.7	11		46.5	4		0.0	328		43.7	27
40.00		58.6	19		26.5	4		29.3	4		58.6	5
50.00		61.8	8		12.0	15		68.2	4		61.8	10
60.00		62.3	11		3.2	11		85.3	5		62.3	7
70.00		62.4	11		0.0	245		90.6	4		62.4	10
80.00		62.5	10		1.7	6		92.1	4		62.5	13
90.00		62.5	13		7.3	60		92.5	4		62.5	13
100.00		62.5	15		15.5	12		92.6	4		62.5	14





Performed splits:


Round 41 split on feature 2 with incentive 5179.85

Round 42 split on feature 1 with incentive 2768.04

Round 43 split on feature 1 with incentive 78.48

Seed: 4059059292





Seed: 4059059292

Optimal pricing strategy: C1(TF): 50.00, C2(TT, FF): 60.00, C3(FT): 60.00

UCB with context generation:

Price	TT,TF,FT,FF	Gaps	Pulls
10.00		100.0	8
20.00		69.4	8
30.00		43.2	8
40.00		22.1	8
50.00		6.7	16
60.00		0.0	274
70.00		12.1	19
80.00		41.3	8
90.00		62.2	8
100.00		71.1	8

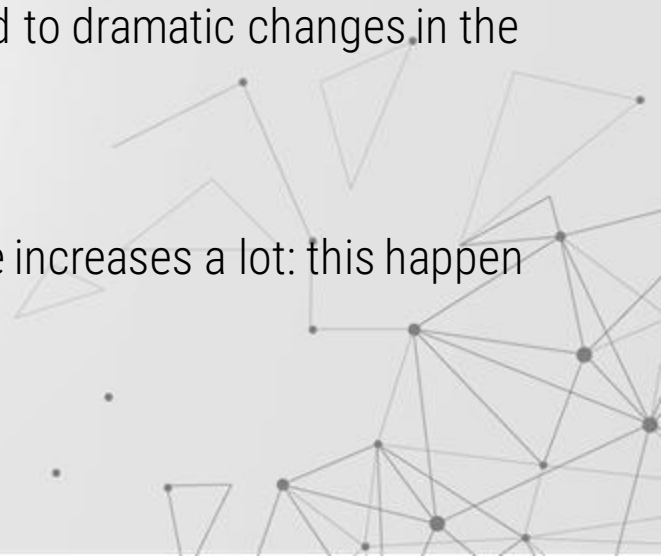
TS with context generation:

Price	TT,TF,FT,FF	Gaps	Pulls
10.00		100.0	8
20.00		69.4	8
30.00		43.2	8
40.00		22.1	8
50.00		6.7	9
60.00		0.0	292
70.00		12.1	8
80.00		41.3	8
90.00		62.2	8
100.00		71.1	8

NO SPLITS FOUND


FINAL CONSIDERATIONS

- We found some discrepancies between the two algorithm, leading to potential differences in the early convergence
- We are estimating a conversion rate, so a small change can lead to dramatic changes in the expected profit
- If a small amount of people buy at a very high price, the revenue increases a lot: this happen when we overestimate the conversion rate for high prices



UCB VS TS

- UCB tends to be optimistic, since it uses upper bounds, hence it gives more shots to higher prices
- TS converges faster, while UCB "wastes" time exploring expensive arms
- Thanks to the nature of the bounds, the regret remains logarithmic, so it converges in a decent number of rounds





Step 5

Online bidding

Without discriminating between customer classes

OUR NUMBERS

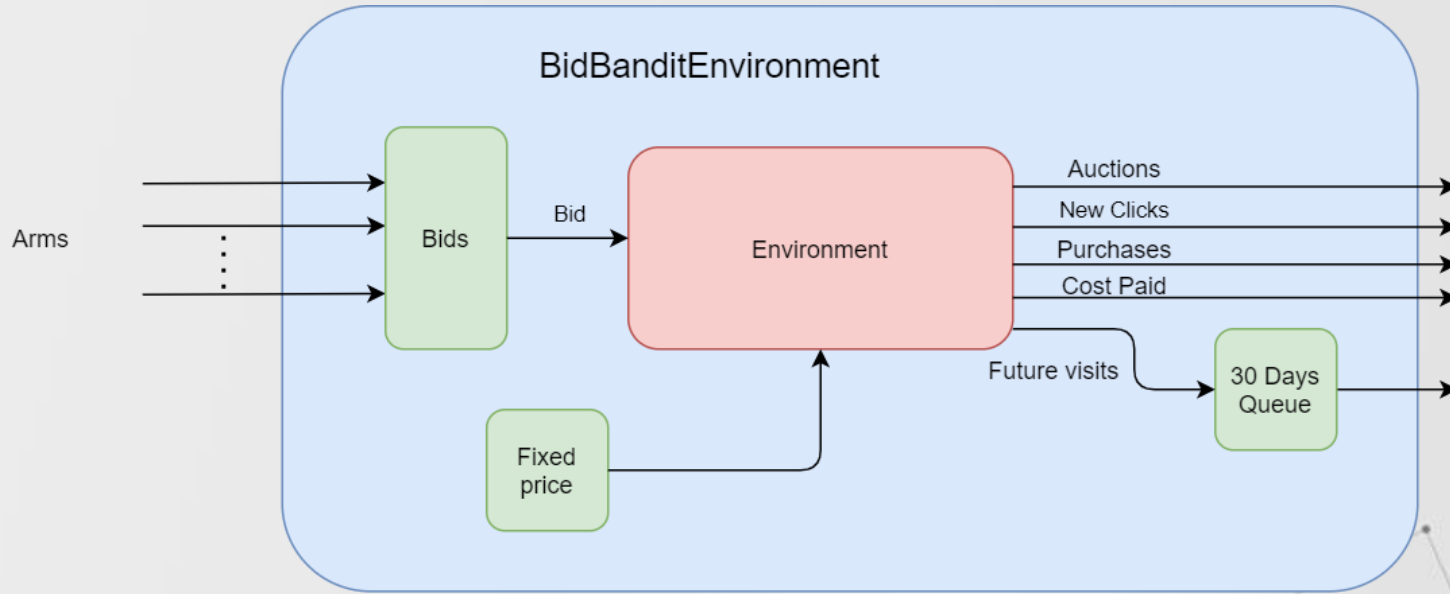
10

BIDS

1

PRICE





- Symmetric to PriceBanditEnvironment.
- Now environment returns also the total number of auctions
- Our random variable is now the probability of winning an auction



SAFETY CONSTRAINT

We don't want to pull arms that could give us negative profit!

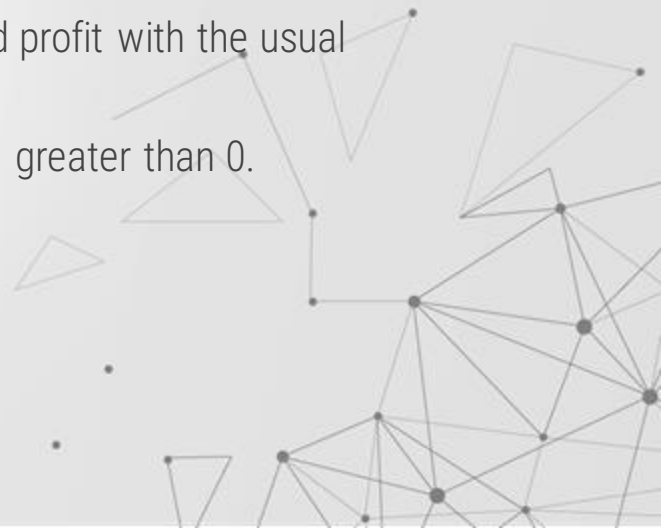
- The main difference between this step and step 3 is the introduction of a safety constraint.
- A *Safe arm* is defined as an arm which revenue won't be negative with a certain probability.
- We use Gaussian regression to approximate the data on our random variable on a normal curve.
- Given the mean and variance of the normal curve, we extract the n -th percentile.

SAFETY CONSTRAINT

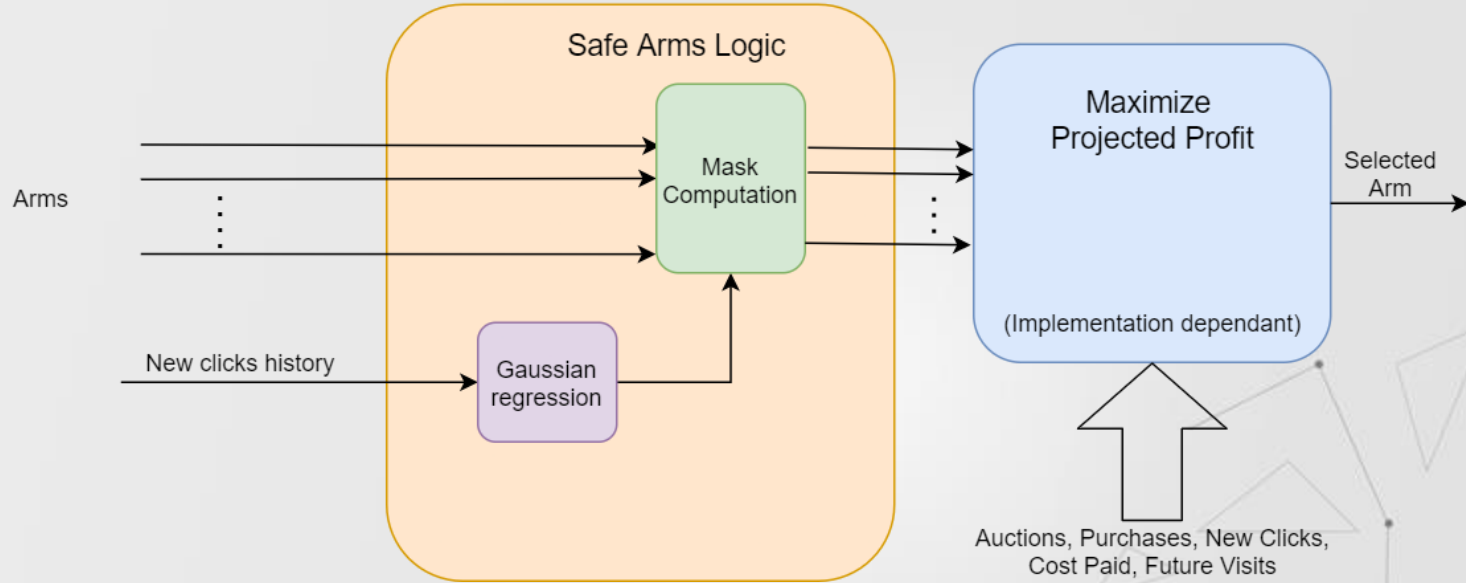
As a result, we have a value for the random variable which is the lowest on $(100-n)\%$ of the cases!

For example if we take the 20th percentile, we are sure that 80% of the times that's the lowest value we will find.

- Given this safety-bound value, we can compute the expected profit with the usual formula.
- An arm is safe if and only if its new expected profit is strictly greater than 0.
- The learner will only pull an arm which is tagged as *safe*.



OptimalBidLearner








$$\frac{\textit{New Clicks}_a}{\textit{Auctions}_a} + \sqrt{\frac{2 \cdot \log(t)}{\textit{Auctions}_a}}$$

- Standard UCB computation
- The denominator is the number of auctions, since we are approximating the probability of winning an auction








Seed: 3077083550





Seed: 538541279





Step 6

Online joint pricing/bidding

Without discriminating between customer classes



OUR NUMBERS

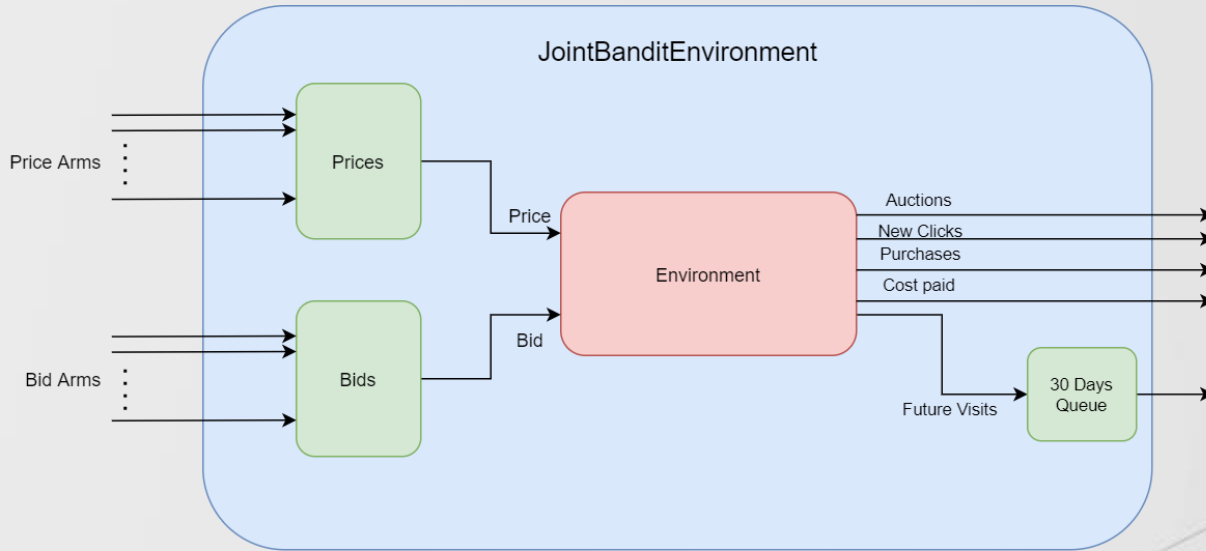
10

PRICES

10

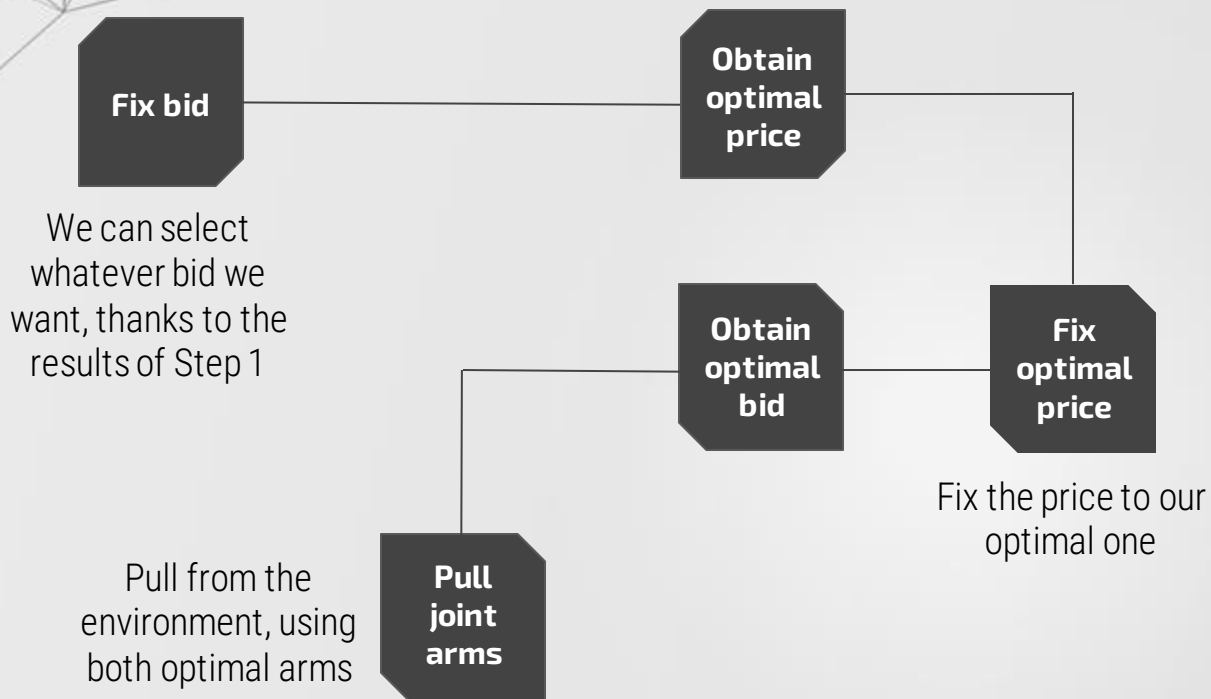
BIDS





- The environment is similar to *PriceBanditEnvironment* and *BidBanditEnvironment*, but incorporates the features of both.
- Accepts both a price and a bid arm.
- No fixed values.

NEXT ARMS CHOICE



Axis projection

- We noticed that every quantity, depends only on one of our set of arms.
- This means that if we are smart, we can re-use all the functionalities from the previous learners.
- If we want to get the average of that quantity for each price arm, we can fix a price and sum together the realizations even if coming from different bids. Of course this remains valid for the specular case when a quantity depends only on the bid.
- This allows us to avoid exploring all the 10×10 possible joint arms.
- All the computations we have to do, can be recycled from *OptimalPriceLearner* and *OptimalBidLearner*, we just need to properly project the data!





QUANTITIES DEPENDENCY

Conversion rate, future visits

**All depend
on the price**

New clicks, auctions, cost per
click

**All depend on the
bid**




Same features as before

- As we already said, the projection technique allowed us to re-use efficiently the code (and the samples).
- The computation of the confidence bound for our random variables is the same as before.
- We still have our safety constraint mechanism which operates only on the bids.



Same results as before

- The convergence results are very similar to the previous cases.
- The absence of cross-dependency allowed our learner to learn as fast as before even though now there are two different random variables to estimate.



Round Robin

- As always, some Round Robin rounds must be performed at the beginning of the algorithm to have some initial data to work with.
- One naïve approach would be to cycle through all the price/bids combinations.
- This would imply to run 100 round robin rounds. Not really smart.








Round Robin

- We can exploit the single arm dependency of all the quantities to avoid pulling unnecessary combinations.
- Then we only have to cycle through 10 prices and 10 bids, which can be done in 10 rounds. Due to the reward delay, actually we have to Round Robin for at least 40 rounds.
- In principle we could select whatever arm we want, the important thing is that after 10 rounds we must have pulled all the prices and all the bids.
- For simplicity we opted to "explore" the diagonal, so pulling (price 1, bid 1), then (price 2, bid 2) and so on.








Seed: 4264432570





Seed: 1966222620






Step 7

Online joint pricing/bidding

Discriminating between customer classes



OUR NUMBERS

10

PRICES

10

BIDS



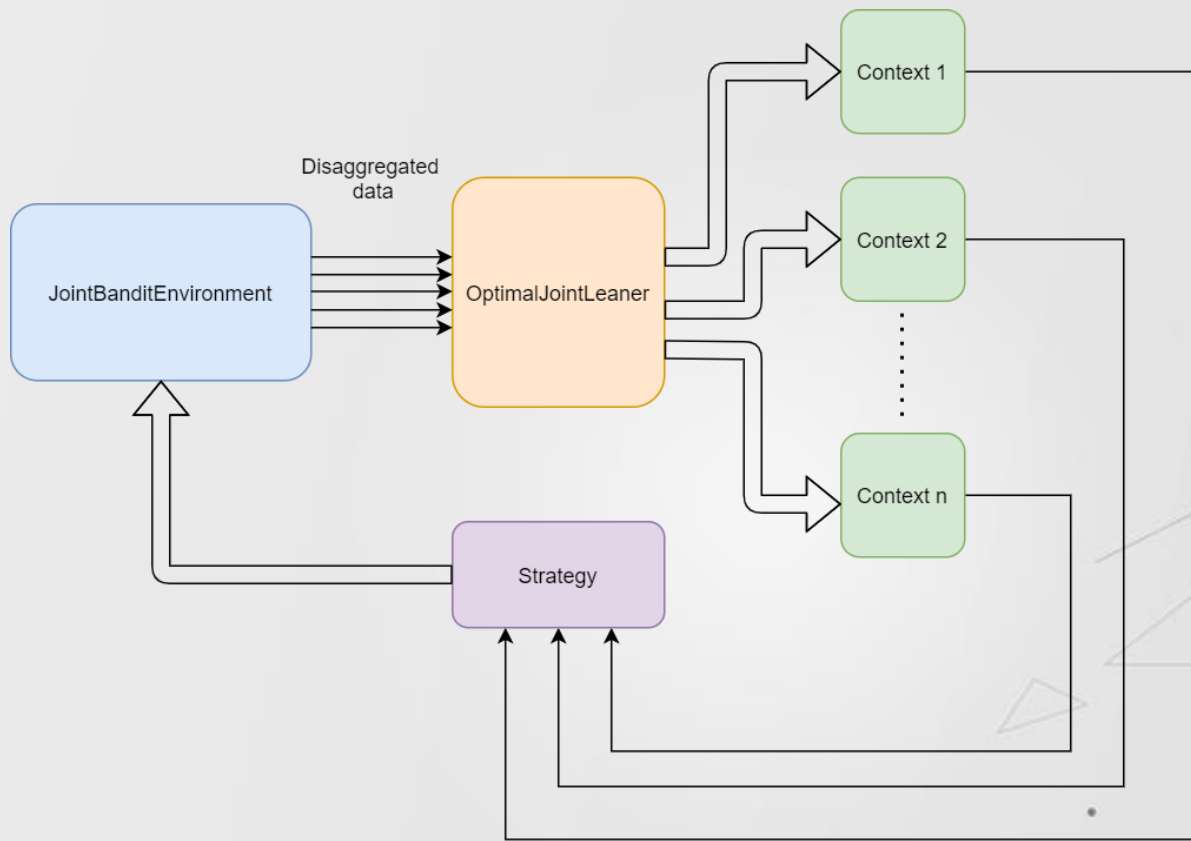
CONTEXTS

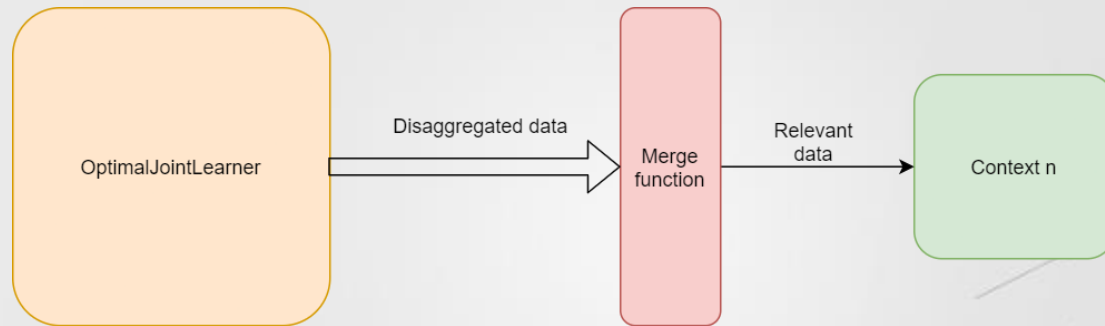
Step 4 finds
suitable contexts

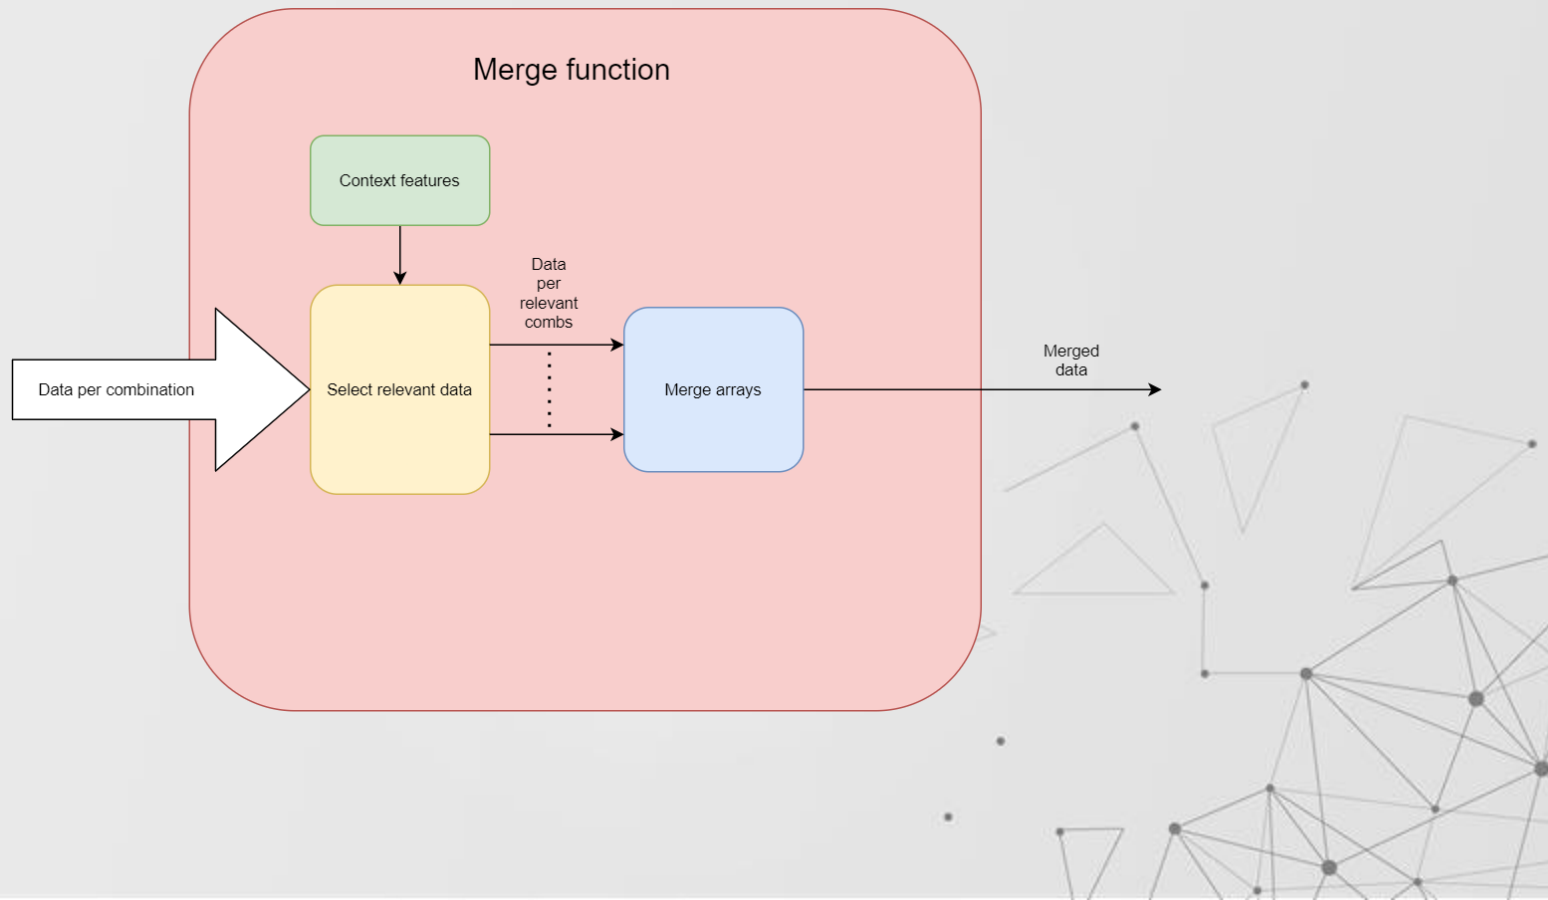


Step 7 uses these
contexts









LEARNING

- Given the merge, all the contexts will run exactly like a small step 6. Indeed we took the very same code from step 6 as after the merge we have aggregated data which matches the features of the context.
- We still choose the arms in two steps.
- There still is the safety constraint.
- Every context provides its best choice and the learner will pack all the choices in a strategy.
- Also our UCB implementation resembles the one of step 6.








Results

- As expected, the results are very similar to the previous steps.
- The overall regret depends also on the results of the step 4 that provided the contexts.








Seed: 3511939391





Seed: 1740212098





The background is a light gray gradient. On the left side, there is a complex network of thin gray lines connecting various black dots of different sizes, forming a web-like structure. Scattered across the entire background are numerous thin, light gray outlines of triangles of various sizes and orientations. Some of these triangles are isolated, while others are part of the network on the left. In the upper right corner, there are small, faint clusters of dots.

The end!